

```
library(ggplot2)
```

## Import Data

```
# Import the dataset
nflx_df <- read.csv("C:/Users/benog/OneDrive/Documents/Grad School/USD/ADS 503 Predictive Modeling/Data/nflx.csv")

# Add headers to the dataframe
#names(glass_df) <- c("Id_number", "RI", "Na", "Mg", "Al", "Si", "K", "Ca", "Ba", "Fe", "Type_of_glass")

head(nflx_df)
```

```
##      Date   Open   High   Low  Close Adj.Close  Volume
## 1 2018-02-05 262.00 267.90 250.03 254.26    254.26 11896100
## 2 2018-02-06 247.70 266.70 245.00 265.72    265.72 12595800
## 3 2018-02-07 266.58 272.45 264.33 264.56    264.56  8981500
## 4 2018-02-08 267.08 267.62 250.00 250.10    250.10  9306700
## 5 2018-02-09 253.85 255.80 236.11 249.47    249.47 16906900
## 6 2018-02-12 252.14 259.15 249.00 257.95    257.95  8534900
```

## Find Null Values

```
# Find number of null values for each attribute
missing_data <- colSums(is.na(nflx_df))
missing_data
```

```
##      Date   Open   High   Low  Close Adj.Close  Volume
##      0         0         0         0         0         0         0
```

## Find Outliers

```
# Use Tukey's method to find outliers
find_outliers <- function(x) {
  q <- quantile(x, probs=c(0.25, 0.75), na.rm=TRUE)
  iqr <- IQR(x, na.rm=TRUE)
  lower_bound <- q[1] - 1.5 * iqr
  upper_bound <- q[2] + 1.5 * iqr
  outliers <- x[x < lower_bound | x > upper_bound]
  return(outliers)
}

# Exclude the Date column
nflx_df_no_date <- nflx_df[, !names(nflx_df) %in% "Date"]

# Find the outliers for each predictor variable
outliers_list <- lapply(nflx_df_no_date, find_outliers)

# Display the outliers
names(outliers_list) <- names(nflx_df_no_date)
outliers_list
```

```

## $Open
## numeric(0)
##
## $High
## numeric(0)
##
## $Low
## numeric(0)
##
## $Close
## numeric(0)
##
## $Adj.Close
## numeric(0)
##
## $Volume
## [1] 18986100 18525800 20369200 18972900 19145500 20307900 33866500 18222800
## [9] 18389900 22490900 22960000 58410400 21746300 18260700 17427300 17183100
## [17] 20156400 32610900 18461000 19039300 19616000 21698800 23685700 20360300
## [25] 21397600 19330100 18620100 19500400 21181200 18871200 26621000 17941400
## [33] 18740200 18054100 31287100 17718000 23832800 38258900 23429900 21730000
## [41] 18200300 17939700 23177600 21084800 21605600 18399000 24499000 24991400
## [49] 20373700 17405700 32637500 22897400 58904300 32346000 24324700 20047500
## [57] 22542300

```

```
# Use Tukey's method to find outliers and output to a dataframe
```

```

get_outliers_df <- function(x, dates) {
  q <- quantile(x, probs=c(0.25, 0.75), na.rm=TRUE)
  iqr <- IQR(x, na.rm=TRUE)
  lower_bound <- q[1] - 1.5 * iqr
  upper_bound <- q[2] + 1.5 * iqr
  outliers <- x[x < lower_bound | x > upper_bound]
  outlier_dates <- dates[x < lower_bound | x > upper_bound]
  return(data.frame(Value = outliers, Date = outlier_dates))
}

```

```
# Add the corresponding date to the outlier's value
```

```

outliers_list <- lapply(nflx_df_no_date, get_outliers_df, dates = nflx_df$Date)

```

```
# Combine outlier values and dates into a single dataframe
```

```

outliers_df <- do.call(rbind, outliers_list)
outliers_df

```

```

##           Value      Date
## Volume.1 18986100 2018-03-05
## Volume.2 18525800 2018-03-06
## Volume.3 20369200 2018-03-12
## Volume.4 18972900 2018-03-28
## Volume.5 19145500 2018-03-29
## Volume.6 20307900 2018-04-16
## Volume.7 33866500 2018-04-17
## Volume.8 18222800 2018-06-13
## Volume.9 18389900 2018-06-21
## Volume.10 22490900 2018-06-25

```

```
## Volume.11 22960000 2018-07-16
## Volume.12 58410400 2018-07-17
## Volume.13 21746300 2018-07-18
## Volume.14 18260700 2018-07-30
## Volume.15 17427300 2018-08-27
## Volume.16 17183100 2018-10-10
## Volume.17 20156400 2018-10-16
## Volume.18 32610900 2018-10-17
## Volume.19 18461000 2018-10-18
## Volume.20 19039300 2018-10-24
## Volume.21 19616000 2018-10-26
## Volume.22 21698800 2018-10-29
## Volume.23 23685700 2018-10-30
## Volume.24 20360300 2018-10-31
## Volume.25 21397600 2018-12-21
## Volume.26 19330100 2019-01-04
## Volume.27 18620100 2019-01-07
## Volume.28 19500400 2019-01-11
## Volume.29 21181200 2019-01-15
## Volume.30 18871200 2019-01-17
## Volume.31 26621000 2019-01-18
## Volume.32 17941400 2019-01-22
## Volume.33 18740200 2019-04-16
## Volume.34 18054100 2019-04-17
## Volume.35 31287100 2019-07-18
## Volume.36 17718000 2019-07-22
## Volume.37 23832800 2019-09-20
## Volume.38 38258900 2019-10-17
## Volume.39 23429900 2019-10-18
## Volume.40 21730000 2020-01-22
## Volume.41 18200300 2020-01-23
## Volume.42 17939700 2020-01-24
## Volume.43 23177600 2020-04-21
## Volume.44 21084800 2020-04-22
## Volume.45 21605600 2020-07-10
## Volume.46 18399000 2020-07-13
## Volume.47 24499000 2020-07-16
## Volume.48 24991400 2020-07-17
## Volume.49 20373700 2020-08-26
## Volume.50 17405700 2020-10-21
## Volume.51 32637500 2021-01-20
## Volume.52 22897400 2021-04-21
## Volume.53 58904300 2022-01-21
## Volume.54 32346000 2022-01-24
## Volume.55 24324700 2022-01-27
## Volume.56 20047500 2022-01-31
## Volume.57 22542300 2022-02-01
```

Create Box & Whisker Plots

```
# Generate a distinct color for each variable
set_colors <- function(index) {
  colors <- c("lightgreen", "lightcoral", "lightpink", "lightyellow", "lightgrey", "lightblue")
  colors[(index %% length(colors)) + 1]
```

```

}

# Create boxplots for each variable in nflx_df except Date
boxplots <- lapply(seq_along(names(nflx_df)), function(i) {
  var <- names(nflx_df)[i]
  if (var != "Date") {
    ggplot(nflx_df, aes_string(x = "factor(1)", y = var)) +
      geom_boxplot(fill = set_colors(i), color = "black") +
      labs(x = "", y = var) +
      ggtitle(paste("Boxplot of", var))
  }
})

```

```

## Warning: 'aes_string()' was deprecated in ggplot2 3.0.0.
## i Please use tidy evaluation idioms with 'aes()'.
## i See also 'vignette("ggplot2-in-packages")' for more information.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.


```


```


# Remove NULL values from the list
boxplots <- Filter(Negate(is.null), boxplots)


# Print each boxplot
for (plot in boxplots) {
  print(plot)
}

```




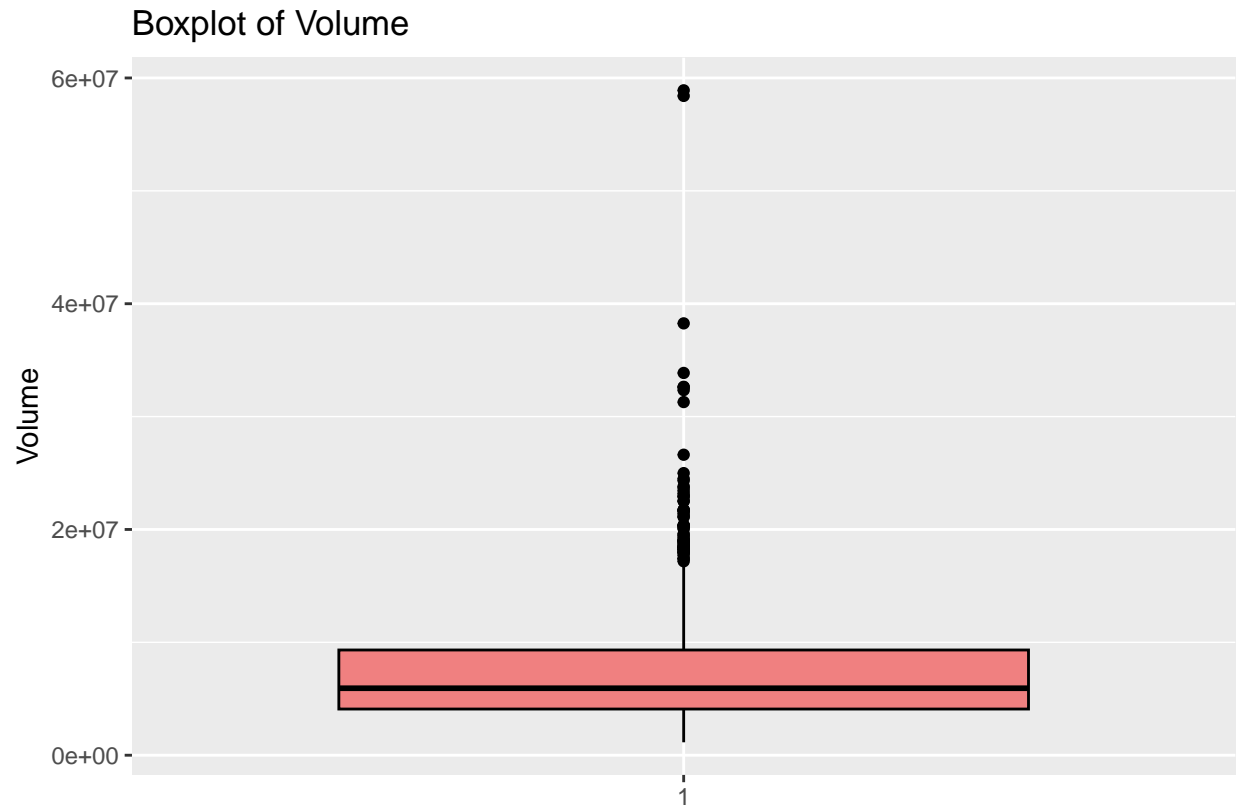










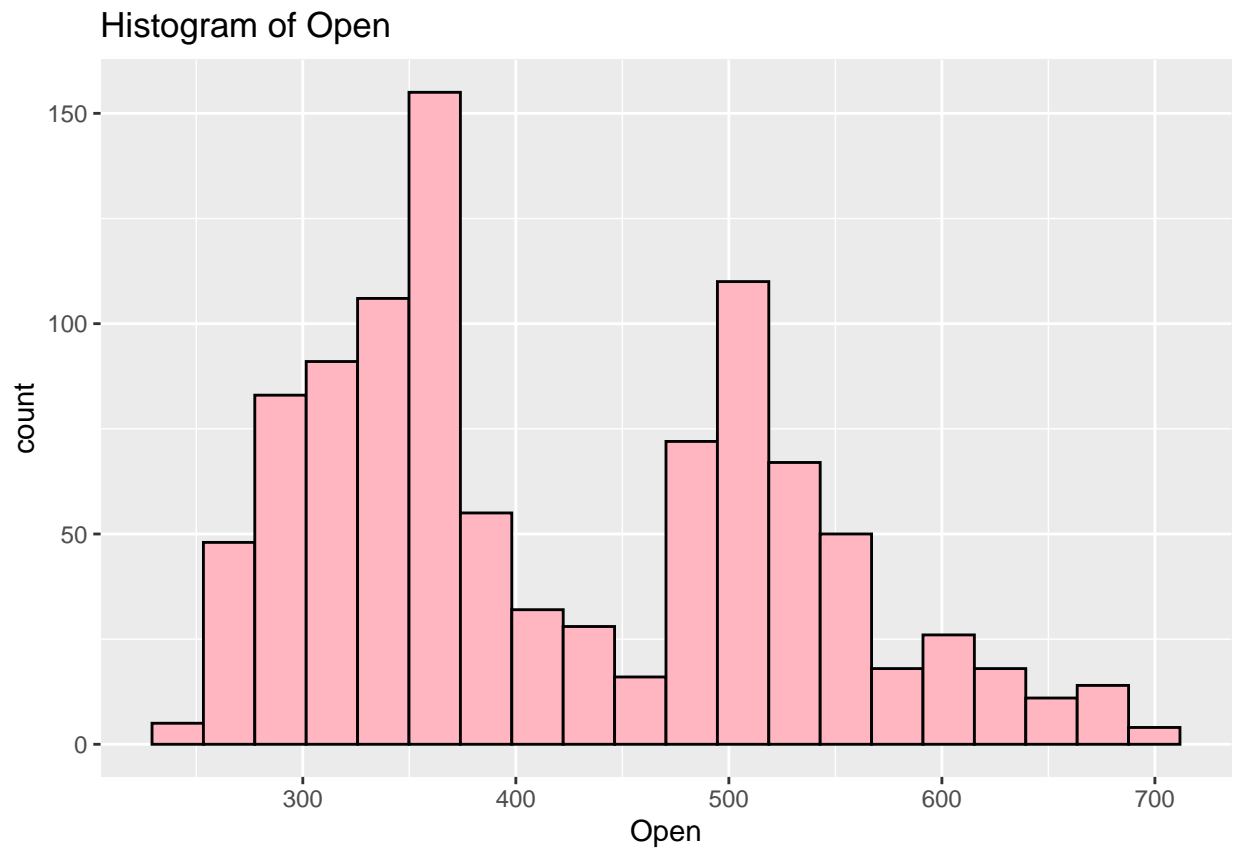



#### Create Histograms

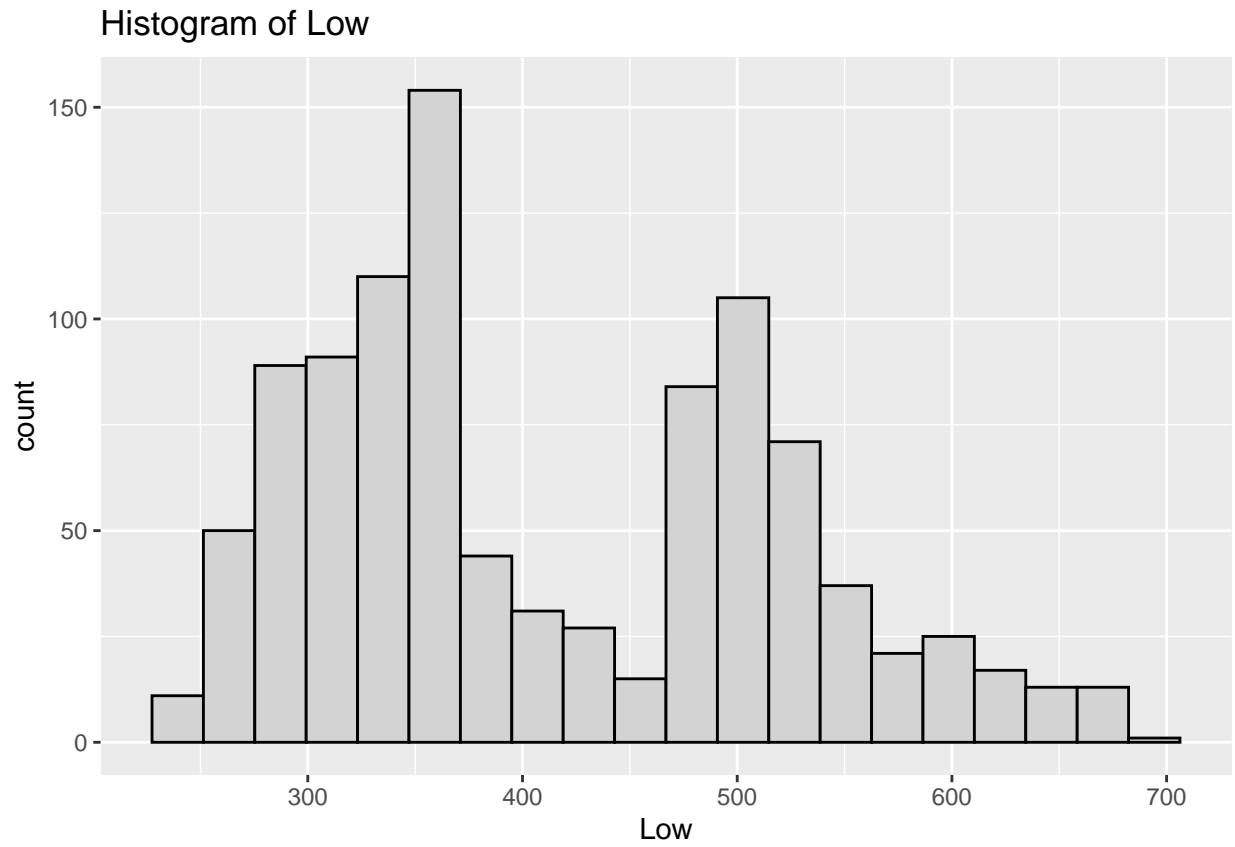
```
# Create histograms for each variable in nflx_df except Date
histograms <- lapply(seq_along(names(nflx_df)), function(i) {
  var <- names(nflx_df)[i]
  if (var != "Date") {
    ggplot(nflx_df, aes_string(x = var)) +
      geom_histogram(bins = 20, fill = set_colors(i), color = "black") +
      labs(x = var) +
      ggtitle(paste("Histogram of", var))
  }
})


# Remove NULL values from the list
histograms <- Filter(Negate(is.null), histograms)

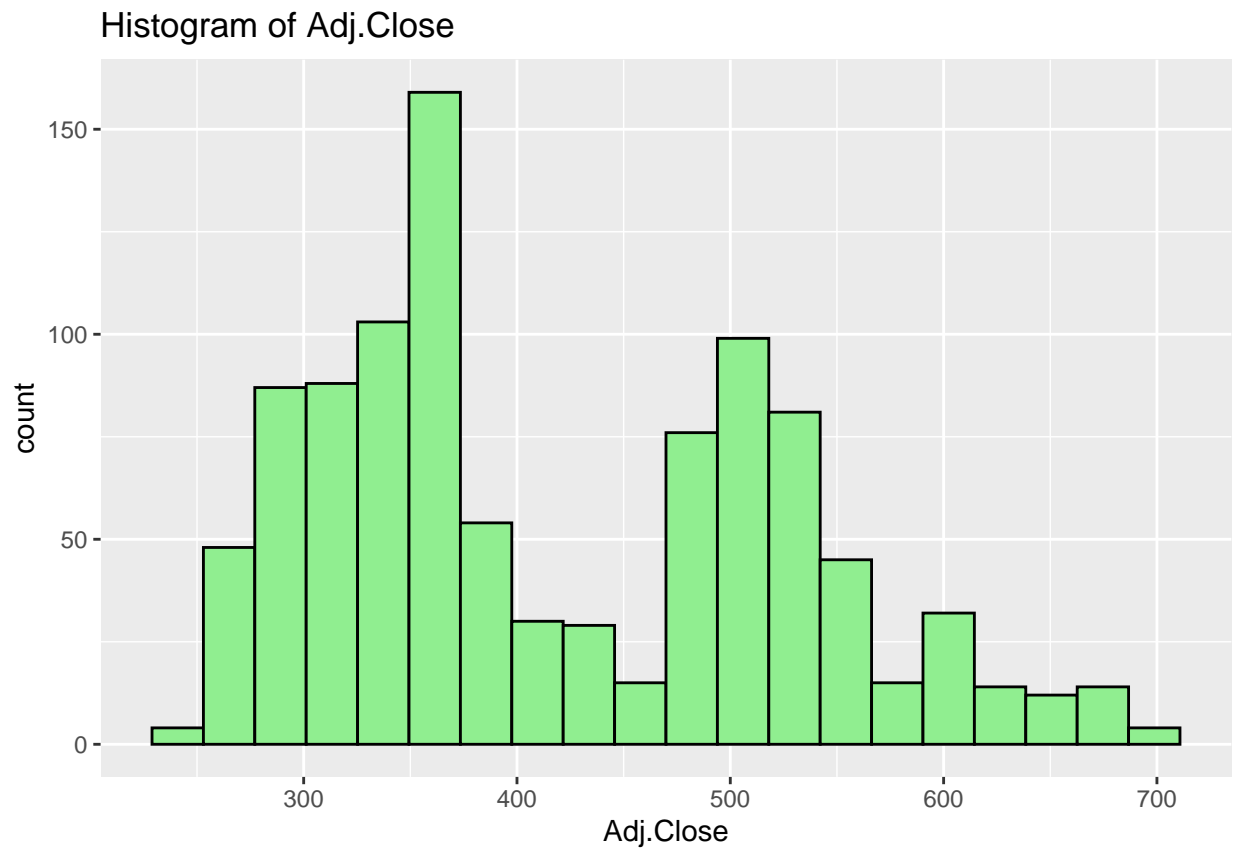
# Print each histogram
for (plot in histograms) {
  print(plot)
}
```












### Histogram of Volume



Create Scatterplots

```
# Generate a distinct color for each variable
set_color_scatter <- function(index) {
  colors <- c("skyblue", "lightcoral", "lightpink", "lightyellow", "lightgrey", "lightblue")
  colors[(index %% length(colors)) + 1]
}

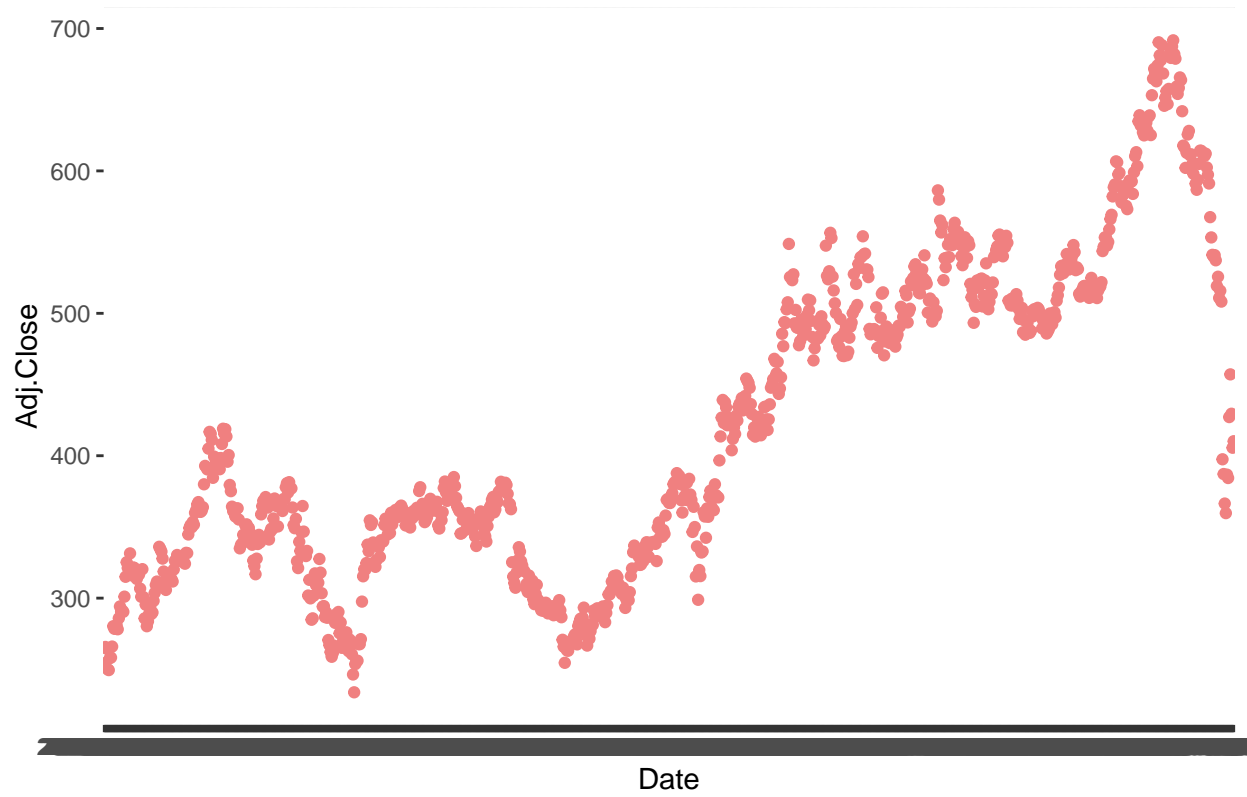
# Create scatterplots for each variable in nflx_df showing the relationship with Adj. Close
scatterplots <- lapply(seq_along(names(nflx_df)), function(i) {
  var <- names(nflx_df)[i]
  if (var != "Adj.Close") {
    ggplot(nflx_df, aes_string(x = var, y = "Adj.Close")) +
      geom_point(color = set_color_scatter(i)) +
      labs(x = var, y = "Adj.Close") +
      ggtitle(paste("Scatterplot of", var, "vs. Adj. Close"))
  }
})

# Remove NULL values from the list
scatterplots <- Filter(Negate(is.null), scatterplots)


# Print each scatterplot
for (plot in scatterplots) {
  print(plot)
}
```



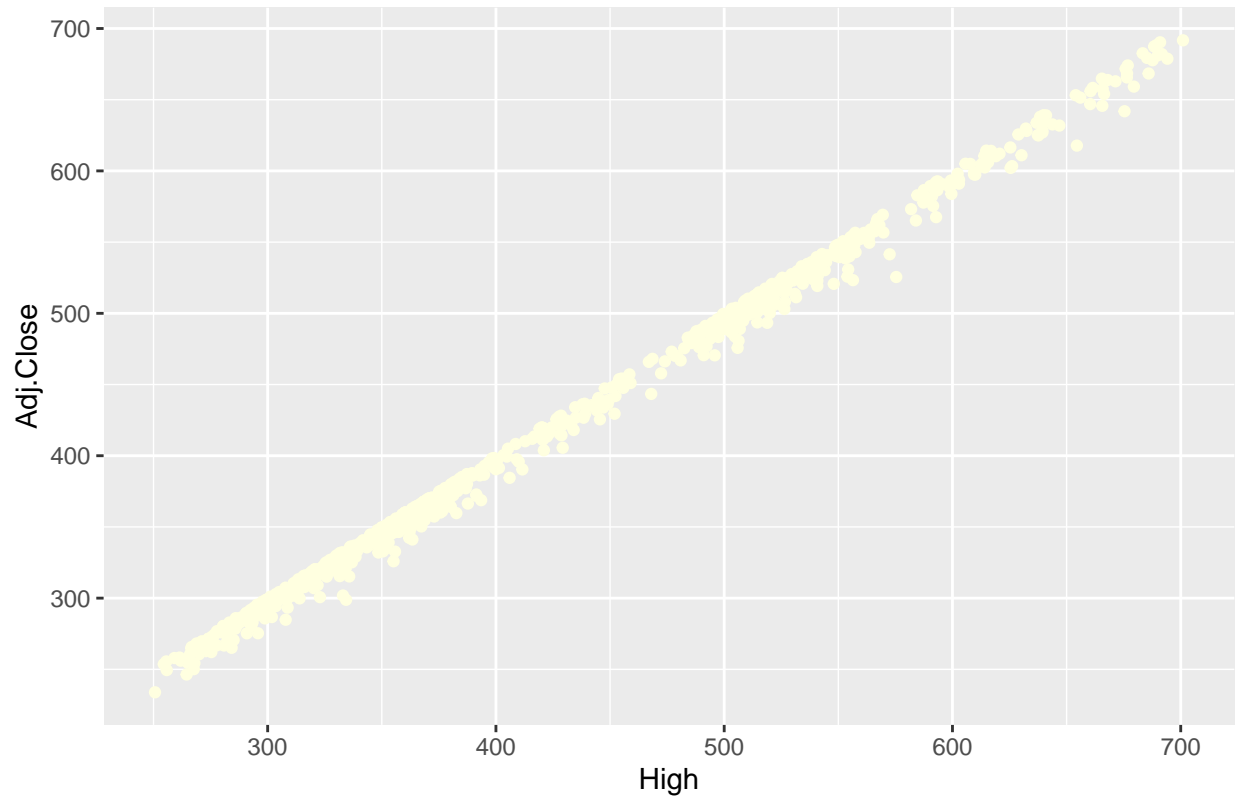
Scatterplot of Date vs. Adj. Close




Scatterplot of Open vs. Adj. Close

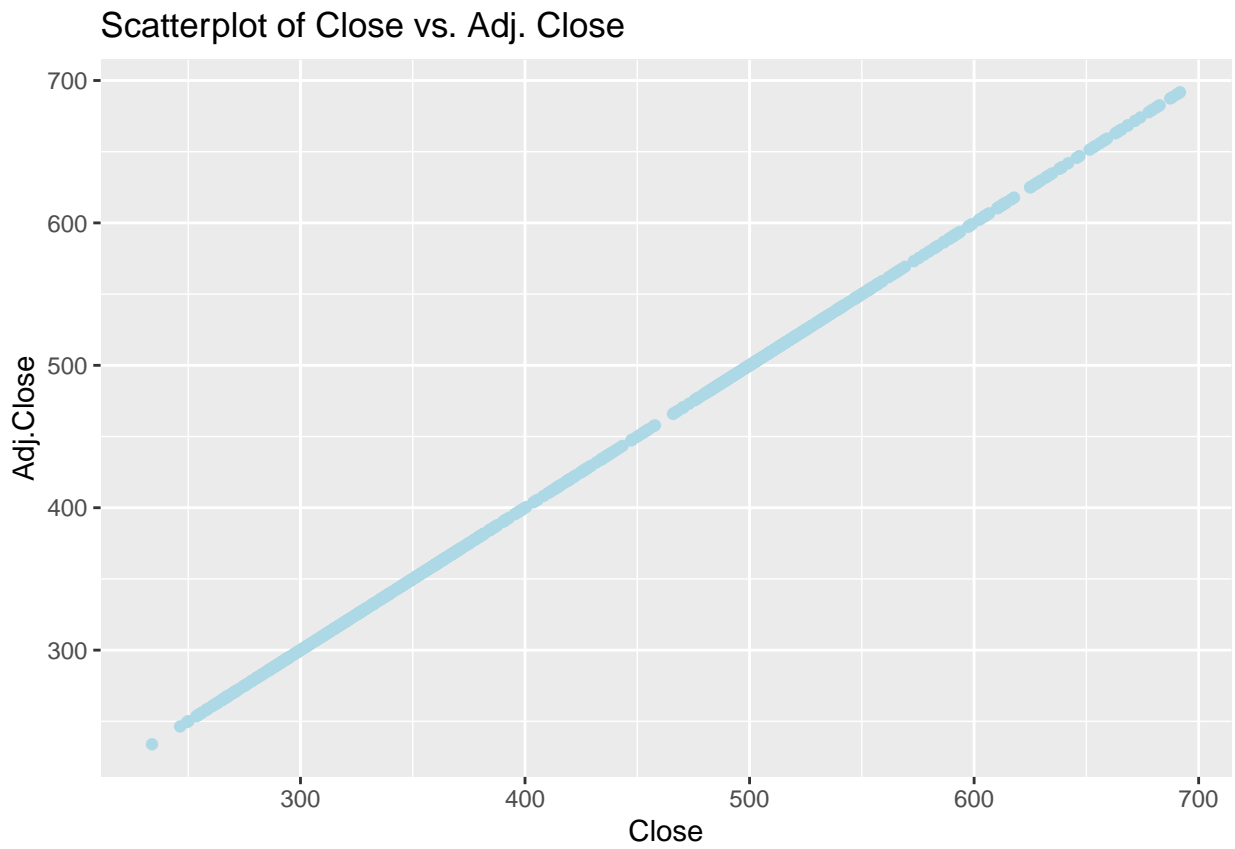


Scatterplot of High vs. Adj. Close




Scatterplot of Low vs. Adj. Close





Scatterplot of Volume vs. Adj. Close



TODO: Not sure if using log transformations or any transformations is good for stock market data. This could lead to distorting true values and predictions since there are actual open, high, low, etc. numbers that represent an actual price or volume in time.

```
# Create a copy of nflx_df
nflx_df_log <- nflx_df

# Exclude the "Date" and "Adj.Price" columns
cols_to_transform <- names(nflx_df_log)[!(names(nflx_df_log) %in% c("Date", "Adj.Price"))]

# Apply log transformation to numeric columns
for(col in cols_to_transform) {
  if(is.numeric(nflx_df_log[[col]])) {
    nflx_df_log[[col]] <- log(nflx_df_log[[col]])
  }
}

# Print the head of the transformed dataframe
print(head(nflx_df_log))
```

##	Date	Open	High	Low	Close	Adj.Close	Volume
## 1	2018-02-05	5.568345	5.590614	5.521581	5.538357	5.538357	16.29172
## 2	2018-02-06	5.512218	5.586124	5.501258	5.582443	5.582443	16.34887
## 3	2018-02-07	5.585674	5.607455	5.577198	5.578068	5.578068	16.01068
## 4	2018-02-08	5.587548	5.589568	5.521461	5.521861	5.521861	16.04625
## 5	2018-02-09	5.536744	5.544396	5.464298	5.519339	5.519339	16.64323


```
## 6 2018-02-12 5.529984 5.557407 5.517453 5.552766 5.552766 15.95967
```

Create Histograms Using Normalized Data


```
# Create histograms for each variable in nflx_df except Date
histograms <- lapply(seq_along(names(nflx_df_log)), function(i) {
  var <- names(nflx_df_log)[i]
  if (var != "Date" && var != "Adj.Close") {
    ggplot(nflx_df_log, aes_string(x = var)) +
      geom_histogram(bins = 20, fill = set_colors(i), color = "black") +
      labs(x = var) +
      ggtitle(paste("Histogram of", var))
  }
})

# Remove NULL values from the list
histograms <- Filter(Negate(is.null), histograms)


# Print each histogram
for (plot in histograms) {
  print(plot)
}
```



Histogram of High







Histogram of Close

